

# DeepSurv in Social Science: Modelling Refugee Journey Duration

---

## Purpose of this Notebook

- This notebook applies *DeepSurv*, a Cox proportional hazards deep neural network, in a social science setting. I study whether *DeepSurv* outperforms the predictive accuracy of traditional Cox models. The basis for comparison is Harrel's c-index, which in this case measures how well a model ranks the arrival times of refugees. The c-index of the best traditional Cox model is equal to 0.72.
- 

## Step 1: Import Packages

- The script first imports required packages. Those include the deep learning package *Lasagne*, the widely-used Python packages *numpy*, *pandas* and *matplotlib* as well as *DeepSurv*. You need to download and install [DeepSurv from Github \(https://github.com/jaredleekatzman/DeepSurv\)](https://github.com/jaredleekatzman/DeepSurv) prior to the analysis.

```
In [1]: import sys
sys.path.append('../deepsurv')
import deep_surv

from deepsurv_logger import DeepSurvLogger, TensorboardLogger
import utils
import viz

import numpy as np
import pandas as pd

import lasagne
import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline
```

```
/Users/timfingerhut/anaconda3/lib/python3.6/site-packages/h5py/__i
nit__.py:36: FutureWarning: Conversion of the second argument of i
ssubdtype from `float` to `np.floating` is deprecated. In future,
it will be treated as `np.float64 == np.dtype(float).type`.
  from ._conv import register_converters as _register_converters
```

## Step 2: Select the Dataset

The following three lines of code

1. load the dataset,
2. transform the data into a *Pandas* Dataframe (the most common data structure in Python, analogous to a csv file with column headers),
3. and output the first lines of the dataset.

The data appears to be in good shape. Please note that continuous variables, such as age, need to be normalized to a 0 to 1 scale prior to employing *DeepSurv* (mmx refers to the minimum-maximum normalization procedure, the 'imp' ending signifies that missing data was imputed).

```
In [2]: train_dataset_fp = './2006_final.csv' # make sure to set the Python
        working directory to the path of your dataset and insert the name o
        f your csv file
        train_df = pd.read_csv(train_dataset_fp)
        train_df.head()
```

Out[2]:

	fail_1	dur	country1	country2	country4	country5	country6	country7	country8	count
0	1	30	1	0	0	0	0	0	0	
1	1	1	1	0	0	0	0	0	0	
2	1	60	1	0	0	0	0	0	0	
3	1	60	1	0	0	0	0	0	0	
4	1	45	1	0	0	0	0	0	0	

5 rows × 67 columns

## Step 3: Prepare the Data for Event History Analysis

- The data is almost ready for *DeepSurv*. As in traditional event history analysis, the *failure* and *duration* variables still need to be indicated (in this case 'fail\_1' and 'dur'). The event / failure indicator is saved as 'e', time 't' is coded in days in this case and the covariates are saved under 'x').

```
In [3]: # event_col is the header in the df that represents the 'Event / Status' indicator
# time_col is the header in the df that represents the event time
def dataframe_to_deepsurv_ds(df, event_col = 'fail_1', time_col = 'dur'):
    # Extract the event and time columns as numpy arrays
    e = df[event_col].values.astype(np.int32)
    t = df[time_col].values.astype(np.float32)

    # Extract the patient's covariates as a numpy array
    x_df = df.drop([event_col, time_col], axis = 1)
    x = x_df.values.astype(np.float32)

    # Return the deep surv dataframe
    return {
        'x' : x,
        'e' : e,
        't' : t
    }

# If the headers of the csv change, you can replace the values of
# 'event_col' and 'time_col' with the names of the new headers
# You can also use this function on your training dataset, validation dataset, and testing dataset
train_data = dataframe_to_deepsurv_ds(train_df, event_col = 'fail_1', time_col = 'dur')
```

## Step 4: Set Hyperparameters

- Based on a manual hyperparameter tuning focused on the *learning rate*, I retain the following hyperparameters. Advanced users can implement a random or Bayesian optimization procedure. For an example of a random hyperparameter search, visit [Jared Lee Katzman's GitHub \(https://github.com/jaredleekatzman/DeepSurv/tree/master/hyperparam\\_search\)](https://github.com/jaredleekatzman/DeepSurv/tree/master/hyperparam_search).

```
In [4]: hyperparams = {
    'L2_reg': 10.0,
    'batch_norm': True,
    'dropout': 0.4,
    'hidden_layers_sizes': [100, 100],
    'learning_rate': 3e-2,
    'lr_decay': 0.001,
    'momentum': 0.9,
    'n_in': train_data['x'].shape[1],
    'standardize': True
}
```

## Step 5: Train `DeepSurv`

- The dataframe and hyperparameters being defined, it is now time to **train `DeepSurv`**. Tensorboard allows you to monitor the training progress in real-time.
- The code below also defines the **update function** (= amsgrad) as well as the **number of training epochs** (= 2000).

```
In [5]: # Create an instance of DeepSurv using the hyperparams defined above
model = deep_surv.DeepSurv(**hyperparams)

# DeepSurv can now leverage TensorBoard to monitor training and validation
# This section of code is optional. If you don't want to use the tensorboard logger
# Uncomment the below line, and comment out the other three lines:
# logger = None

experiment_name = 'test_experiment_tim'
logdir = './logs/tensorboard/'
logger = TensorboardLogger(experiment_name, logdir=logdir)

# Now we train the model
update_fn=lasagne.updates.amsgrad # The type of optimizer to use. \
                                   # Check out http://lasagne.readthedocs.io/en/latest/modules/updates.html \
                                   # for other optimizers
                                   # to use
n_epochs = 2000

# If you have validation data, you can add it as the second parameter to the function
metrics = model.train(train_data, n_epochs=n_epochs, logger=logger,
update_fn=update_fn)
```

```
2019-07-30 13:53:45,675 - Training step 0/2000      |
| - loss: 23.3311 - ci: 0.6639
2019-07-30 13:54:43,876 - Training step 250/2000  |***
| - loss: 6.3844 - ci: 0.8368
2019-07-30 13:55:44,226 - Training step 500/2000  |*****
| - loss: 6.2731 - ci: 0.8734
2019-07-30 13:56:42,873 - Training step 750/2000  |*****
| - loss: 6.2329 - ci: 0.8915
2019-07-30 13:57:41,429 - Training step 1000/2000 |*****
| - loss: 6.1705 - ci: 0.8971
2019-07-30 13:58:42,531 - Training step 1250/2000 |*****
| - loss: 6.1805 - ci: 0.9046
2019-07-30 13:59:41,270 - Training step 1500/2000 |*****
***      | - loss: 6.1290 - ci: 0.9070
2019-07-30 14:00:39,638 - Training step 1750/2000 |*****
*****   | - loss: 6.1346 - ci: 0.9047
2019-07-30 14:01:37,559 - Finished Training with 2000 iterations in 472.25s
```

## Step 6: Output and visualize the Results

- Finally, it is time to print and visualize *DeepSurv*'s metrics.
- Compared to the traditional Cox model (c-index = 0.72), *DeepSurv* orders refugees' arrival times far more accurately (c-index = 0.9102).

```
In [6]: # Print the final metrics
print('Train C-Index:', metrics['c-index'][-1])
# print('Valid C-Index: ', metrics['valid_c-index'][-1])

# Plot the training / validation curves
viz.plot_log(metrics)
```

Train C-Index: (1999, 0.9101721615528376)

